

Advanced Machine Learning

Normalizing Flow

Alexandre Vérine

Master Intelligence Artificielle, Systèmes, Données

18/02/2022



Program for today

- 1 Normalizing Flow in theory
 - How does it work ?
 - What can we do with it ?
 - How is it trained ?
- 2 Structure of Normalizing flow
 - Theoretical Structures
 - First Structures
 - Coupling - Glow
 - Residual Network - ResFlow
- 3 Current Limits of Normalizing Flow
- 4 Practical Work with Python

A Normalizing Flow is usually seen as:

- a *generative model*,
- a *bijective mapping*,
- an *invertible neural network*,
- a *density estimator*.




Figure: A mapping between two probability distributions
Point to point




Figure: A mapping between two probability distributions
Subset to subset

Normalizing Flow

A Normalizing Flow is a bijective function between a data space \mathcal{X} and a latent space \mathcal{Z} , both subset of \mathbb{R}^d .

$$\begin{aligned} F : \mathcal{X} &\longmapsto \mathcal{Z} \\ x &\longmapsto z = F(x) \end{aligned}$$

Data and Latent Distributions

In theory, a NF maps a target distribution P , ie the data distribution to a simple latent distribution Q .

Usually, Q is set to be a Normal Gaussian multivariate distribution $\mathcal{N}(0_d, I_d)$. p and q are respectively the probability densities of P and Q .

How does it work ?

In practice, the mapping is *not perfect*. P^* induces a distribution Q and similarly, the latent distribution Q induces \hat{P} , which is the learned distribution. The forward pass F is called the *Normalizing* direction while the inverse pass F^{-1} is called the *Generative* direction.




Figure: 1D Normalizing Flow process.

Induced Probabilities ?

Change of Variable Formula

For a bijective and continuous fonction F and a latent distribution Q , the distribution induced by Q and F is defined through the *change of variable formula*:

$$\forall \mathbf{x} \in \mathcal{X}, \quad \hat{p}(\mathbf{x}) = |\det \text{Jac}_F(\mathbf{x})| q(F(\mathbf{x})). \quad (1)$$

Induced Probabilities ?

$$\forall \mathbf{x} \in \mathcal{X}, \quad \hat{p}(\mathbf{x}) = |\det \text{Jac}_F(\mathbf{x})| q(F(\mathbf{x})).$$




Figure: Example of 1D mapping

Induced Probabilities ?

$$\forall \mathbf{x} \in \mathcal{X}, \quad \hat{p}(\mathbf{x}) = |\det \text{Jac}_F(\mathbf{x})| q(F(\mathbf{x})).$$




Figure: Example of 1D mapping

Induced Probabilities ?

$$\forall \mathbf{x} \in \mathcal{X}, \quad \hat{p}(\mathbf{x}) = |\det \text{Jac}_F(\mathbf{x})| q(F(\mathbf{x})).$$




Figure: Example of 1D mapping

Induced Probabilities ?

$$\forall \mathbf{x} \in \mathcal{X}, \quad \hat{p}(\mathbf{x}) = |\det \text{Jac}_F(\mathbf{x})| q(F(\mathbf{x})).$$




Figure: Example of 1D mapping

Induced Probabilities ?

$$\forall \mathbf{x} \in \mathcal{X}, \quad \hat{p}(\mathbf{x}) = |\det \text{Jac}_F(\mathbf{x})| q(F(\mathbf{x})).$$




Figure: Example of 1D mapping

Density Estimation

To perform density estimation:

- 1 Draw $\mathbf{x} \sim P^*$,
- 2 Compute $F(\mathbf{x})$ and $|\det \text{Jac}_F(\mathbf{x})|$,
- 3 Compute $\hat{p}(\mathbf{x}) = q(F(\mathbf{x}))|\det \text{Jac}_F(\mathbf{x})|$.




Figure: 1D Normalizing Process of Density Estimation.

To perform data generation:

- 1 Draw $\mathbf{z} \sim Q$,
- 2 Compute $\mathbf{x} = F^{-1}(\mathbf{z})$.




Figure: 1D Normalizing Flow process of Generation.

How is it trained ? - Generative models

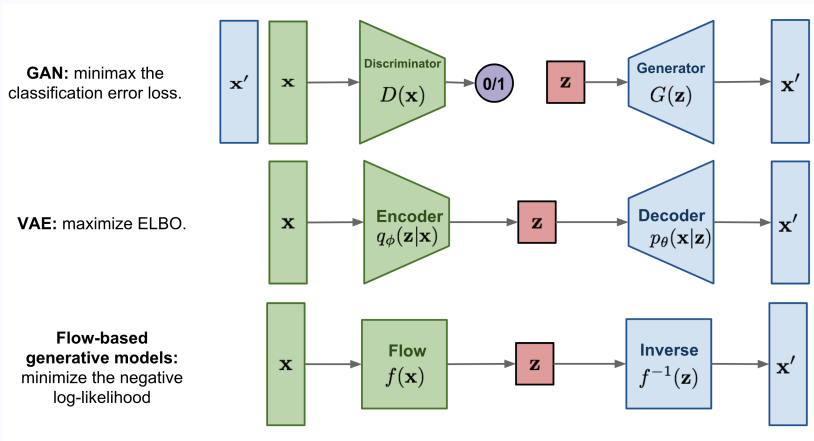


Figure: Different generative models [4]

Loss

The objective is to approximate P^* with \hat{P} . We can minimize the Kullback-Leiber Divergence :

$$\theta = \arg \min_{\theta} \mathcal{D}_{\text{KL}}(P^* \parallel \hat{P}).$$

This is equivalent to maximizing the log likelihood :

$$\theta = \arg \max_{\theta} \mathbb{E}_{x \sim \mathcal{X}} [\log \hat{p}(x)] .$$

$$\mathcal{D}_{\text{KL}}(P^* \parallel \hat{P}) = \int_{\mathcal{X}} p^*(\mathbf{x}) \log \left(\frac{p^*(\mathbf{x})}{\hat{p}(\mathbf{x})} \right) d\mathbf{x}$$

$$\text{nll} = -\mathbb{E}_{\mathbf{x} \sim \mathcal{X}} [\log \hat{p}(\mathbf{x})]$$

Learning steps

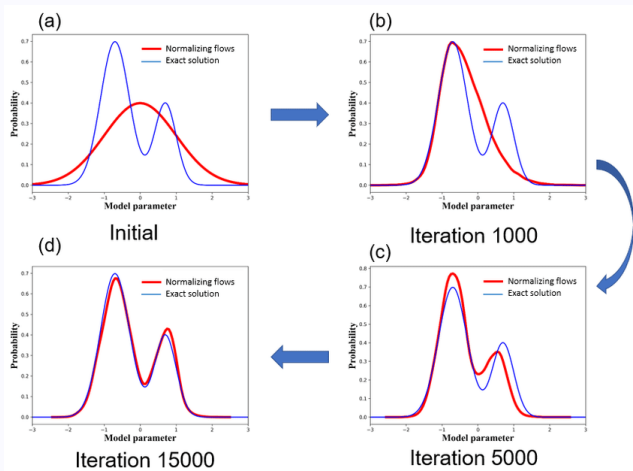


Figure: Learning Process for a 1D Normalizing Flow [5].

The entropy can be used as an intuitive understanding of the loss. We have :

$$H(P^*, \hat{P}) = \mathcal{D}_{\text{KL}}(P^* \| \hat{P}) + H(P^*) \geq H(P^*).$$

In practice, the entropy per dimension is used to compare different models on a same dataset :

$$\text{bpd} = \frac{\sum_{\mathbf{x} \in \mathcal{X}} \log(\hat{p}(\mathbf{x})/2^K)}{d} = \frac{\text{nll}}{dK \log(2)}$$

where K is the number of different bits to encode the data. For instance, we usually have $256 = 2^8$ different values per pixel per channel, thus $K = 8$.

Normalizing Flow

A Normalizing Flow is a Neural Network for which we can efficiently compute :

- The forward pass F , ie the normalizing direction,
- the inverse pass F^{-1} , ie the generative direction,
- the determinant of the Jacobian matrix J_{acF} for every \mathbf{x} .

A normalizing Flow is a Neural Network

Usually, a Neural Net is represented as the combination of atomic function $f_i : \mathbf{x} \mapsto \sigma(\mathbf{W}_i \mathbf{x} + \mathbf{b}_i)$ such that $F = f_n \circ \dots \circ f_0$.




Figure: Representation of a Neural Network.

A normalizing Flow is a Neural Network

For a Normalizing Flow, every atomic block f_i satisfies the properties of F , i.e. bijective and the Jacobian must be efficiently computable.




Figure: Atomic representation of the Normalizing Flow.

Therefore, we have:

$$\log \det \text{Jac}_F(\mathbf{x}) = \sum_{i=0}^n \log \det \text{Jac}_{f_i}(\mathbf{x}_i)$$

A normalizing Flow is a Neural Network




Figure: Representation of an Invertible Neural Network.

Elementwise Flow

The element wise flows are based on elementwise bijective non linear scalar function. Let $h_i : \mathbb{R} \mapsto \mathbb{R}$ be the scalar value bijection :

$$f_i(x_i^1, \dots, x_i^d) = \left(h(x_i^1), \dots, h(x_i^d) \right)^T$$

Determinant of the Jacobian matrix

Here, $\text{Jac}_{f_i}(\mathbf{x}) = \text{Diag} \left(\frac{dh}{dx} \Big|_{x=x_1}, \dots, \frac{dh}{dx} \Big|_{x=x_d} \right)$, so
 $|\det \text{Jac}_{f_i}(\mathbf{x})| = \prod_i^d \frac{dh}{dx} \Big|_{x=x_i}$.

Limits

- No mixing of variables

Linear Flows

Let \mathbf{A}_i be an invertible matrix of $\mathbb{R}^{d \times d}$ and \mathbf{b}_i a vector of \mathbb{R}^d . A Linear Flow is based on the linear operation :

$$f_i(\mathbf{x}) = \mathbf{A}_i \mathbf{x} + \mathbf{b}_i$$

Determinant of the Jacobian matrix

Here, $\text{Jac}_{f_i}(\mathbf{x}) = \mathbf{A}_i$, so $|\det \text{Jac}_{f_i}(\mathbf{x})| = |\det \mathbf{A}_i| \neq 0$.

Limits

Poor Expressiveness.

Planar Flows

Let $h : \mathbb{R} \mapsto \mathbb{R}$, h' its derivative, \mathbf{u} and \mathbf{w} two vectors of \mathbb{R}^d and a scalar b . The Planar flow is based on atomic blocks f_i such that:

$$f_i(\mathbf{x}) = \mathbf{x} + \mathbf{u}h(\mathbf{w}^T\mathbf{x} + b)$$

Determinant of the Jacobian matrix

Here, $\text{Jac}_{f_i}(\mathbf{x}) = \mathbf{I}_d + \mathbf{u}h'(\mathbf{w}^T\mathbf{x} + b)\mathbf{w}^T$, so
 $|\det \text{Jac}_{f_i}(\mathbf{x})| = 1 + h'(\mathbf{w}^T\mathbf{x} + b)\mathbf{u}^T\mathbf{w}$.

Limits

No closed form of the inverse.

Radial Flows

Let \mathbf{x}_0 a vector of \mathbb{R}^d , α and β two scalars. The Radial Flow is based on atomic blocks f_i such that:

$$f_i(\mathbf{x}) = \mathbf{x} + \frac{\beta}{\alpha + \|\mathbf{x} - \mathbf{x}_0\|} (\mathbf{x} - \mathbf{x}_0)$$

Determinant of the Jacobian matrix

Here, $|\det \text{Jac}_{f_i}(\mathbf{x})| = \left[1 + \frac{\beta}{\alpha + \|\mathbf{x} - \mathbf{x}_0\|}\right]^{d-1} \left[1 + \frac{\beta}{\alpha + \|\mathbf{x} - \mathbf{x}_0\|} + \frac{\|\mathbf{x} - \mathbf{x}_0\|}{(\alpha + \|\mathbf{x} - \mathbf{x}_0\|)^2}\right]$.

Limits

No closed form of the inverse.

Radial and Planar Flows




Figure: Example of Planar and Radial Flows [2].

Coupling Function

A coupling function is coupling disjoint partitions of the dataset. For an input $\mathbf{x} \in \mathbb{R}^d$, we consider a split $(\mathbf{x}^A, \mathbf{x}^B) \in \mathbb{R}^k \times \mathbb{R}^{d-k}$. Usually, $k = d/2$. We consider a bijective function, generally elementwise, $\mathbf{h} : \mathbf{x} \mapsto \mathbf{h}(\mathbf{x}, \theta) \in \mathbb{R}^k$ parameterized by θ . We can train a small neural network such that θ becomes a function of \mathbf{x}^A or \mathbf{x}^B . We can therefore define the atomic coupling function :

$$f_i\left(\begin{bmatrix} \mathbf{x}^A \\ \mathbf{x}^B \end{bmatrix}\right) = \begin{bmatrix} \mathbf{y}^A \\ \mathbf{y}^B \end{bmatrix} = \begin{bmatrix} \mathbf{h}(\mathbf{x}^A, g(\mathbf{x}^B)) \\ \mathbf{x}^B \end{bmatrix}.$$

The reverse is :

$$\mathbf{f}_i^{-1}\left(\begin{bmatrix} \mathbf{y}^A \\ \mathbf{y}^B \end{bmatrix}\right) = \begin{bmatrix} \mathbf{x}^A \\ \mathbf{x}^B \end{bmatrix} = \begin{bmatrix} \mathbf{h}^{-1}(\mathbf{y}^B, g(\mathbf{x}^B)) \\ \mathbf{x}^B \end{bmatrix}.$$

Coupling Function




Figure: Atomic coupling block.

Coupling Function




Figure: Inverse Atomic coupling block.

Coupling Function




Figure: Composed Atomic coupling block.

Determinant of the Jacobian Matrix

Here, $\text{Jac}_{f_i}(\mathbf{x}) = \begin{pmatrix} \frac{\partial h}{\partial \mathbf{x}^A} |_{\mathbf{x}^A, g(\mathbf{x}^B)} & \frac{\partial h}{\partial \mathbf{x}^B} |_{\mathbf{x}^A, g(\mathbf{x}^B)} \\ \mathbf{0} & \mathbf{I}_k \end{pmatrix}$, thus

$$|\det \text{Jac}_{f_i}(\mathbf{x})| = \left| \frac{\partial h}{\partial \mathbf{x}^A}(\mathbf{x}^A) \right|$$




Figure: Atomic coupling block.

Additive Coupling

Additive Coupling Flow

The additive coupling flow has atomic function defined with an elementwise function $h_i : \mathbb{R} \mapsto \mathbb{R}$ and the scalar function $g_i : \mathbb{R}^d \mapsto \mathbb{R}$:

$$h_i(x_i) = x_i + g_i(\mathbf{x}^B).$$

Therefore,

$$f_i \left(\begin{bmatrix} \mathbf{x}^A \\ \mathbf{x}^B \end{bmatrix} \right) = \begin{bmatrix} \mathbf{x}^A + g(\mathbf{x}^B) \\ \mathbf{x}^B \end{bmatrix}.$$

Determinant of the Jacobian matrix

Here, $|\det \text{Jac}_{f_i}| = 1$. The flow is said to be volume preserving.

Limits

Very deep structure and poor expressiveness.

Affine Coupling

Affine Coupling Flow

The affine coupling flow has atomic function defined with an elementwise function $h_i : \mathbb{R} \mapsto \mathbb{R}$ and the scalar function $g_i : \mathbb{R}^d \mapsto \mathbb{R}$:

$$h_i(x_i) = g_i^1(\mathbf{x}^B)x_i + g_i^2(\mathbf{x}^B).$$

Therefore,

$$f_i \left(\begin{bmatrix} \mathbf{x}^A \\ \mathbf{x}^B \end{bmatrix} \right) = \begin{bmatrix} g^1(\mathbf{x}^B)\mathbf{x}_A + g^2(\mathbf{x}^B) \\ \mathbf{x}^B \end{bmatrix}.$$

Determinant of the Jacobian matrix

Here, $|\det \text{Jac}_{f_i}| = (g_i^1(\mathbf{x}^B))^k$.

Limits

Very deep structure.

Residual Flow

A residual flow is composed of residual atomic blocks f_i defined with function $\mathbf{g}_i : \mathbb{R}^d \mapsto \mathbb{R}^d$ such that:

$$f_i(\mathbf{x}) = \mathbf{x}_i + \mathbf{g}(\mathbf{x}_i).$$

In \mathbf{g} , we can stack Convolutional layers, batchnorms, activations, etc...

Determinant of the Jacobian Matrix

Here, the determinant of the log determinant of the Jacobian matrix is estimated by a Russian Roulette Estimator, and a Skilling-Hutchingson trace estimator [1]:

$$\log \det \text{Jac}_{f_i}(\mathbf{v}\mathbf{x}) = \mathbb{E}_{n \sim \text{Geo}(p), \mathbf{v} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[\sum_{k=1}^n \frac{(-1)^{k+1}}{k} \frac{\mathbf{v}^T [\text{Jac}_g(\mathbf{x})^k \mathbf{v}]}{p(n \geq k)} \right]$$

Inverse

A residual Network is invertible only if $Lip(g_i) < 1$. We can compute the inverse with the Inverse-point algorithm:

Algorithm 1 Inverse via fixed-point iteration.

```
 $\mathbf{x}_{i-1} \leftarrow \mathbf{x}_i$   
while NotConverged do  
     $\mathbf{x}_{i-1} \leftarrow \mathbf{x}_i - \mathbf{g}(\mathbf{x}_{i-1})$   
end while
```

Limits

Iterative inverse and approximated Jacobian.

Since Normalizing Flows are diffeomorphism, there are bi-Lipschitz :

$$\forall \mathbf{x}_1, \mathbf{x}_2 \in \mathcal{X}, \quad \|F(\mathbf{x}_1) - F(\mathbf{x}_2)\|_2 \leq L_1 \|\mathbf{x}_1 - \mathbf{x}_2\|_2,$$

and

$$\forall \mathbf{z}_1, \mathbf{z}_2 \in \mathcal{Z}, \quad \|F^{-1}(\mathbf{z}_1) - F^{-1}(\mathbf{z}_2)\|_2 \leq L_2 \|\mathbf{z}_1 - \mathbf{z}_2\|_2.$$

Therefore, some pathological target distributions can be highlighted [3].

Expressivity of a Normalizing Flow : Dense Subset

Let F be L_1 -Lipschitz. Then:

$$\text{TV}(P^*, \hat{P}) \geq \sup_{R, \mathbf{x}_0} \left(P^*(B_{R, \mathbf{x}_0}) - \frac{\gamma\left(\frac{d}{2}, \frac{L_1^2 R^2}{2}\right)}{\Gamma\left(\frac{d}{2}\right)} \right)$$

Therefore, if we find a ball for which the true measure satisfies $P^*(B_{R, \mathbf{x}_0}) > \gamma(\frac{d}{2}, \frac{L_1^2 R^2}{2})/\Gamma(\frac{d}{2})$, then the TV is necessarily strictly positive.




Figure: Example of a pathological target distribution: the subset B_R concentrates most of the weight in $P^*(B_R)$, but $\hat{P}(B_R) = Q(F(B_R))$ can only be as large as $Q(B_{RL_1})$.

Expressivity of a Normalizing Flow : Sparse Subset

Let F^{-1} be L_2 -Lipschitz. We consider the balls centered on $F^{-1}(0)$, we have the lower bound:

$$\text{TV}(P^*, \hat{P}) \geq \sup_R \left(\frac{\gamma\left(\frac{d}{2}, \frac{R^2}{2L_2^2}\right)}{\Gamma\left(\frac{d}{2}\right)} - P^*(B_{R, F^{-1}(0)}) \right)$$

Therefore, if we find a ball for which the the true measure satisfies $P^*(B_{R, F^{-1}(0)}) < \frac{\gamma(d/2, R^2/2L_2^2)}{\Gamma(d/2)}$, then the TV is necessarily strictly positive.




Figure: Example of a pathological target distribution: the subset B_R concentrates little weight in $P^*(B_R)$, but $\hat{P}(B_R) = Q(F(B_R))$ can only be as small as $Q(B_R/L_2)$.

Conclusion

- [1] Ricky T. Q. Chen, Jens Behrmann, David Duvenaud, and Jörn-Henrik Jacobsen. Residual Flows for Invertible Generative Modeling. In *33rd Conference on Neural Information Processing Systems (NeurIPS 2019), Vancouver, Canada.*, July 2020. arXiv: 1906.02735.
- [2] Danilo Jimenez Rezende and Shakir Mohamed. Variational Inference with Normalizing Flows. In *arXiv:1505.05770 [cs, stat]*, June 2016.
- [3] Alexandre Verine, Yann Chevaleyre, Fabrice Rossi, and Benjamin Negrevergne. On the expressivity of bi-Lipschitz normalizing flows. In *ICML Workshop on Invertible Neural Networks, Normalizing Flows, and Explicit Likelihood Models*, June 2021.
- [4] Lilian Weng. Flow-based deep generative models. *lilianweng.github.io/lil-log*, 2018.

- [5] Xuebin Zhao, Andrew Curtis, and Xin Zhang. Bayesian Variational Seismic Tomography using Normalizing Flows. In *EGU General Assembly Conference Abstracts*, EGU General Assembly Conference Abstracts, pages EGU21–1455, April 2021.